

GODENA SIDDARTHA

AI Software Developer | AI Engineer (LLM / GenAI Systems)

+91 8686155152 | gsiddartha117@gmail.com | [linkedin.com/in/siddartha-g-080b23271](https://www.linkedin.com/in/siddartha-g-080b23271) | github.com/SiddarthaGodena-AI

Python • FastAPI • LLMs • Gemini • AWS Bedrock • RAG • LangChain • Multi-Agent Systems • MCP • AWS • GCP / Vertex AI • Docker

PROFESSIONAL SUMMARY

AI Software Developer / AI Engineer with 1+ year of experience building AI backend systems under senior technical mentorship — FastAPI microservices, LLM integration, multi-agent systems, and RAG pipelines with vector search. Independently implemented 3 production AI systems at Inoday, processing 10,000+ customer interactions and 500+ hours of multimedia content. Also building applied computer vision and voice AI automation as self-directed practice.

TECHNICAL SKILLS

Languages: Python, SQL, C++

Backend Development: FastAPI, REST API Design, Microservices Architecture, Asynchronous Programming, Serverless, AWS Lambda, Event-Driven Architecture, Pydantic

AI / ML Integration: LLMs, Gemini, AWS Bedrock, HuggingFace Transformers, RAG Pipelines, FAISS, ChromaDB, LangChain, MCP, Sentence Transformers, Multi-Agent Systems, Prompt Engineering

Computer Vision: OpenCV, YOLO, YOLOv4-tiny, Real-Time Object Detection, Video Analytics, Image Classification

Automation / Voice AI: n8n, Retell AI, Twilio, Webhook-based Event Automation

Cloud Platforms: AWS, S3, Lambda, Transcribe, Bedrock, GCP, Vertex AI, GCS, Gemini API

DevOps & MLOps: Docker, MLflow, Git, CI/CD, GitHub Actions

Data & ML: Pandas, NumPy, scikit-learn, PyTorch, NLP, Sentiment Analysis, OCR

Databases: MySQL, MongoDB, FAISS, ChromaDB

EXPERIENCE

Junior Software Developer — Inoday | Jun 2024 - Aug 2025

Reported directly to a Senior Technical Lead and the company's Founding Partner.

- Independently implemented 3 backend AI systems using FastAPI + Docker, processing unstructured multimedia data (audio, video, documents) end-to-end within project scope set by senior technical leadership.
- Built serverless ingestion workflows on AWS (S3 + Lambda event triggers) that automatically processed 10,000+ customer support interactions without manual intervention.
- Built a multi-agent LLM routing layer with LangChain + Gemini API to dispatch queries by domain instead of a single general-purpose agent.

Artificial Intelligence Intern — Slash Mark | Mar 2024 - Jun 2024

- Built end-to-end ML pipelines for data preprocessing, model training, and evaluation; delivered structured analytical insights to support product decisions.

INDEPENDENT PROJECTS & SELF-DIRECTED DEVELOPMENT

Aug 2025 - Present

CV Object Detection & Analytics API | Aug 2026

Tech Stack: Python, OpenCV, YOLOv4-tiny, FastAPI, Docker, pytest | github.com/SiddarthaGodena-AI/cv-object-detection-api

- Built a REST API for real-time object detection and video analytics using OpenCV's DNN module with a pretrained YOLOv4-tiny model, avoiding a full PyTorch dependency to keep container size and cold-start latency low.
- Implemented image detection, annotated-image, and video-summary-analytics endpoints; validated against public benchmark footage, correctly tracking person/car/truck classes across a 795-frame test video. Covered with automated pytest and containerized with Docker.

Voice AI Automation System — n8n, Retell AI, Twilio | 2025 - 2026

- Architected conversational voice-agent workflows connecting Twilio telephony to Retell AI's real-time voice/LLM engine, with n8n handling orchestration and business logic between calls.
- Implemented multi-step call-handling flows covering intent classification, call routing, and fallback logic, triggered by webhook-based events across the integrated API stack.

PRODUCTION SYSTEMS — INODAY

Jun 2024 - Aug 2025, under senior technical mentorship. These 3 ran in production.

AI Call Sentry — Serverless Customer QA Platform (AWS)

Python, AWS S3/Transcribe/Lambda, FastAPI, Docker, NLP | github.com/SiddarthaGodena-AI/AI-Call-Sentry

- Built a fully serverless pipeline (S3 upload trigger -> Lambda -> Transcribe -> sentiment model) so incoming support calls were processed automatically with no dedicated server to manage.
- Benchmarked sentiment output against manually reviewed call samples before rollout, reaching 88-92% agreement with human review; my mentor used this to sign off on replacing manual QA for the bulk of incoming calls (10,000+ processed).

Talk-to-Data — Multimodal AI Analytics Platform

Python, FastAPI, Vertex AI (Gemini), FAISS, NLP, Computer Vision APIs

- Built REST APIs for natural-language querying across 500+ hours of combined video, audio, image, and text data, so non-technical users could ask questions instead of manually reviewing recordings.
- Designed automated speaker-analytics and insight-extraction pipelines using Vertex AI (Gemini) to summarize session content without a human reviewing each recording first.

Voice-Driven Multi-Agent Troubleshooting System (GenAI + RAG)

Python, FastAPI, Vertex AI, FAISS, LangChain, Gemini API

- Built a multi-agent API routing incoming queries to domain-specialized LLM agents via LangChain instead of one general-purpose agent handling everything.
- Designed a FAISS-backed RAG pipeline so agent responses were grounded in retrieved documents rather than the model's unaided output.

ADDITIONAL PRACTICE PROJECTS — INODAY TENURE

Personal builds alongside day-to-day work; not deployed to production.

SmartNews AI — FastAPI news aggregation service with real-time fetch, metadata enrichment, and LLM-generated summaries.

github.com/SiddarthaGodena-AI/AI-NewsUpdates

AI-HRBot — FastAPI + AWS Bedrock + ChromaDB backend for semantic resume-to-job-description matching and ranked candidate output. github.com/SiddarthaGodena-AI/AI-HRBot

EDUCATION

B.Tech in Computer Science Engineering (AI & ML) — Dayananda Sagar University | 2020 - 2024

Higher Secondary (10+2) — Narayana Junior College | 2018 - 2020

CERTIFICATIONS

Elements of AI (University of Helsinki) • Python for Machine Learning and Data Science (Udemy) • Google Cloud Masterclass (Udemy)